



Contemporary C++ in HPC codes

Jakub Gałęcki^{1,2}

October 17th 2023, UQ Computers & Fluids Discussion Group

¹ Department of Aerodynamics, Warsaw University of Technology

² Interdisciplinary Centre for Mathematical and Computational Modelling,
University of Warsaw

Why C++?

- Lingua franca for HPC
- C interop
- Zero cost abstractions
- Rich ecosystem of existing (legacy) software
- Still growing & evolving!

Why C++?

- Lingua franca for HPC
- C interop
- Zero cost abstractions
- Rich ecosystem of existing (legacy) software
- Still growing & evolving!

Why should I care if I'm coding in X?

- What facilities does X provide for the problems of interest?
- How can I improve X?

Why C++?

- Lingua franca for HPC
- C interop
- Zero cost abstractions
- Rich ecosystem of existing (legacy) software
- Still growing & evolving!

Why should I care if I'm coding in X?

- What facilities does X provide for the problems of interest?
- How can I improve X?

Maintainability carries real-world cost

Why C++?

- Lingua franca for HPC
- C interop
- Zero cost abstractions
- Rich ecosystem of existing (legacy) software
- Still growing & evolving!

Why should I care if I'm coding in X?

- What facilities does X provide for the problems of interest?
- How can I improve X?

Maintainability carries real-world cost

Naming: modern C++ = {11, 14}, contemporary C++ ≥ 17

- Background & motivation
- Examples:
 - Miscellanea
 - MPI wrapper
 - Physics injection

- Background & motivation
- Examples:
 - Miscellanea
 - MPI wrapper
 - Physics injection

Please interrupt if you have questions/comments

- Background & motivation
- Examples:
 - Miscellanea
 - MPI wrapper
 - Physics injection

Please interrupt if you have questions/comments

Click to see code: 

The least-squares finite element method, L3STER

Consider the abstract boundary-value problem:

$$\begin{cases} \mathcal{A}(\mathbf{u}) = \mathbf{f} & \text{in } \Omega \\ \mathcal{B}(\mathbf{u}) = \mathbf{g} & \text{on } \partial\Omega \end{cases}$$

We define the least-squares functional:

$$\mathcal{J}(\mathbf{u}; \mathbf{f}, \mathbf{g}) = \frac{1}{2} \left(\|\mathcal{A}(\mathbf{u}) - \mathbf{f}\|_{\Omega,0}^2 + \|\mathcal{B}(\mathbf{u}) - \mathbf{g}\|_{\partial\Omega,0}^2 \right)$$

The least-squares principle

Find $\mathbf{u} \in \mathcal{V}$ such that

$$\forall \tilde{\mathbf{u}} \in \mathcal{V} \quad \mathcal{J}(\mathbf{u}; \mathbf{f}, \mathbf{g}) \leq \mathcal{J}(\tilde{\mathbf{u}}; \mathbf{f}, \mathbf{g})$$

Restrict \mathcal{V} to a finite-dimensional $\mathcal{V}_h = \text{span} \{ \Phi_1, \dots, \Phi_N \}$

$$u = \sum_{i=1}^N u_i \Phi_i, \text{ where } u_i \in \mathbb{R}$$

We now have a convex algebraic minimization problem, where the stationary point must satisfy

$$\begin{aligned} & \left[\int_{\Omega} (\mathcal{A}^T \Phi_i) \mathcal{A} \Phi_j + \int_{\partial\Omega} (\mathcal{B}^T \Phi_i) \mathcal{B} \Phi_j \right] [u_j] \\ &= \left[\int_{\Omega} (\mathcal{A}^T \Phi_i) f + \int_{\partial\Omega} (\mathcal{B}^T \Phi_i) g \right] \end{aligned}$$

Alternative view: weak formulation where the derivatives of the result of the operator w.r.t. the unknowns serve as the trial functions

The domain contribution to the local matrix is given by

$$\mathbf{K} = \left[\int_{\Omega} (\mathcal{A}^T \Phi_i) \mathcal{A} \Phi_j \right] \approx \sum_{q=1}^Q w_q \left[(\mathcal{A}^T \Phi_i(x_q)) \mathcal{A} \Phi_j(x_q) \right] = \mathbf{A}^T \mathbf{W}_d \mathbf{A}$$

- ✓ There's no weak formulation!
- ✓ BCs are trivial to impose (replace \mathcal{A} with \mathcal{B})
- ✓ The system is SPD regardless of \mathcal{A}
- ✓ Error norm is built in
- ✗ \mathcal{A} must be first-order and linear
- ✗ Quadrature order is usually higher than for the Galerkin method, e.g. for 3D convection-diffusion $Q = 8p^3$

Linear first-order operators can be expressed as

$$\mathcal{A} = A_0 + \sum_{i=1}^{dim} A_i \frac{\partial}{\partial x_i}$$

where $A_i : \Omega \times (0, T] \rightarrow \mathbb{R}^{E \times U}$, E being the number of equations and U the number of unknowns (E and U may not be equal).

All we need to define a set of physics is to provide these matrices!

Different physics in different domains can be handled automatically.

- Least-Squares Scalable Spectral Element framework
- Written in C++ 20
- General PDE solver – no equations included in the core
- High premium placed on passing parameters (mostly kernel dimensions) at compile-time
- Low entry barrier: equations + mesh = solution

Miscellaneous C++ bits for HPC

Span = pointer + size

Non-owning view over a contiguous memory region

Convenient API:

- Compile-time or dynamically sized
- Sub-spans
- Range constructor

Span = pointer + size

Non-owning view over a contiguous memory region

Convenient API:

- Compile-time or dynamically sized
- Sub-spans
- Range constructor

```
1 auto v          = std::vector{10, 9, 8, 7, 6, 5, 4, 3, 2, 1};
2 auto v_span     = std::span{v};
3 auto drop3      = v_span.subspan(3);
4 auto drop3_take6 = drop3.subspan(0, 6);
5 std::print("{}\n", drop3_take6);
6 std::sort(drop3_take6.begin(), drop3_take6.end());
7 std::print("{}\n", v);
```



Span = pointer + size

Non-owning view over a contiguous memory region

Convenient API:

- Compile-time or dynamically sized
- Sub-spans
- Range constructor

```
1 auto v          = std::vector{10, 9, 8, 7, 6, 5, 4, 3, 2, 1};  
2 auto v_span     = std::span{v};  
3 auto drop3      = v_span.subspan(3);  
4 auto drop3_take6 = drop3.subspan(0, 6);  
5 std::print("{}\n", drop3_take6);  
6 std::sort(drop3_take6.begin(), drop3_take6.end());  
7 std::print("{}\n", v);
```



[7, 6, 5, 4, 3, 2]

[10, 9, 8, 2, 3, 4, 5, 6, 7, 1]

- Multi-dimensional span
- Non-owning view over a contiguous memory region, with multidimensional access information encoded
- API inspired by Kokkos
- Allows strided access and various memory layouts
- Can serve as common interface for matrix types from different libraries
- Finally multi-argument **operator[]!**
- Sorry, still no support on Godbolt :(

```
1 template<typename T>  
2 class atomic_ref;
```

Atomic reference type

Objects can be safely concurrently accessed from multiple threads, provided all accesses are via `std::atomic_ref`

Supports the usual set of operations and memory orders

Now full support for floating-point types!

```
1 void atomic_increment(double& val, double inc) {  
2     std::atomic_ref{val}.fetch_add(inc /*, mem_order */);  
3 }
```

```
1 void atomic_increment(double& val, double inc) {  
2     std::atomic_ref{val}.fetch_add(inc /*, mem_order */);  
3 }
```

```
1 atomic_increment(double&, double):  
2     mov     rax, qword ptr [rdi]  
3 .LBB0_1:  
4     movq    xmm1, rax  
5     addsd   xmm1, xmm0  
6     movq    rcx, xmm1  
7     lock    cmpxchg qword ptr [rdi], rcx  
8     jne     .LBB0_1  
9     ret
```



```
1 void atomic_increment(double& val, double inc) {  
2     std::atomic_ref{val}.fetch_add(inc /*, mem_order */);  
3 }
```

```
1 atomic_increment(double&, double):  
2     mov     rax, qword ptr [rdi]  
3 .LBB0_1:  
4     movq    xmm1, rax  
5     addsd   xmm1, xmm0  
6     movq    rcx, xmm1  
7     lock    cmpxchg qword ptr [rdi], rcx  
8     jne     .LBB0_1  
9     ret
```



Note the use of CTAD...

MPI in contemporary C++

What you stand to gain:

What you stand to gain:

- Communicator ownership semantics:
 - Own
 - View
 - Duplicate

What you stand to gain:

- Communicator ownership semantics:
 - Own
 - View
 - Duplicate
- Comm buffer type deduction – no more `void*`

What you stand to gain:

- Communicator ownership semantics:
 - Own
 - View
 - Duplicate
- Comm buffer type deduction – no more `void*`
- RAI for MPI requests (and other things)

What you stand to gain:

- Communicator ownership semantics:
 - Own
 - View
 - Duplicate
- Comm buffer type deduction – no more **void***
- RAI for MPI requests (and other things)
- Request tied with payload to automate memory management

What you stand to gain:

- Communicator ownership semantics:
 - Own
 - View
 - Duplicate
- Comm buffer type deduction – no more **void***
- RAI for MPI requests (and other things)
- Request tied with payload to automate memory management
- **Range-enabled MPI calls**

What you stand to gain:

- Communicator ownership semantics:
 - Own
 - View
 - Duplicate
- Comm buffer type deduction – no more **void***
- RAI for MPI requests (and other things)
- Request tied with payload to automate memory management
- **Range-enabled MPI calls**

What you stand to lose:

What you stand to gain:

- Communicator ownership semantics:
 - Own
 - View
 - Duplicate
- Comm buffer type deduction – no more **void***
- RAI for MPI requests (and other things)
- Request tied with payload to automate memory management
- **Range-enabled MPI calls**

What you stand to lose:

- API surface

Range:

- represents an iterable sequence
- exposes iterator/sentinel pair via `.begin()` / `.end()`

Range:

- represents an iterable sequence
- exposes iterator/sentinel pair via `.begin()` / `.end()`

View:

- view over part or all of a range
- possibly with additional transformations applied on top
- standard views are lazily evaluated
- composable via pipeline operator

Range:

- represents an iterable sequence
- exposes iterator/sentinel pair via `.begin()` / `.end()`

View:

- view over part or all of a range
- possibly with additional transformations applied on top
- standard views are lazily evaluated
- composable via pipeline operator

Borrowed range:

- range which does not own the underlying sequence
- iterators to a borrowed range can safely outlive their parent range
- a range type can be queried to check whether it's borrowed

Range:

- represents an iterable sequence
- exposes iterator/sentinel pair via `.begin()` / `.end()`

View:

- view over part or all of a range
- possibly with additional transformations applied on top
- standard views are lazily evaluated
- composable via pipeline operator

Borrowed range:

- range which does not own the underlying sequence
- iterators to a borrowed range can safely outlive their parent range
- a range type can be queried to check whether it's borrowed

Ranges facilities are transparent to the compiler – no run-time penalty

Range:

- represents an iterable sequence
- exposes iterator/sentinel pair via `.begin()` / `.end()`

View:

- view over part or all of a range
- possibly with additional transformations applied on top
- standard views are lazily evaluated
- composable via pipeline operator

Borrowed range:

- range which does not own the underlying sequence
- iterators to a borrowed range can safely outlive their parent range
- a range type can be queried to check whether it's borrowed

Ranges facilities are transparent to the compiler – no run-time penalty

Slideware note: `namespace sr = std::ranges, namespace sv = std::views`

```
1 std::vector v{1, 3, 314, 7, 9, 42};  
2 sr::sort(v);  
3 auto even = v | sv::filter([](int i) { return i%2==0; });  
4 std::print("Sorted even elements: {}\n", even);
```



```
1 std::vector v{1, 3, 314, 7, 9, 42};  
2 sr::sort(v);  
3 auto even = v | sv::filter([](int i) { return i%2==0; });  
4 std::print("Sorted even elements: {}\n", even);
```



Sorted even elements: [42, 314]

```
1 std::vector v{1, 3, 314, 7, 9, 42};  
2 sr::sort(v);  
3 auto even = v | sv::filter([](int i) { return i%2==0; });  
4 std::print("Sorted even elements: {}\n", even);
```



Sorted even elements: [42, 314]

```
1 auto maxit = sr::max_element(std::vector{1, 2, 3, 4, 5});  
2 std::print("Max element: {}\n", *maxit);
```




```
1 std::vector v{1, 3, 314, 7, 9, 42};  
2 sr::sort(v);  
3 auto even = v | sv::filter([](int i) { return i%2==0; });  
4 std::print("Sorted even elements: {}\n", even);
```



Sorted even elements: [42, 314]

```
1 auto maxit = sr::max_element(std::vector{1, 2, 3, 4, 5});  
2 std::print("Max element: {}\n", *maxit);
```



error: no match for 'operator*'
(operand type is 'std::ranges::dangling')

The arguments of various MPI calls get simplified:

- Communicator \rightarrow base object
- (pointer, size, type) \rightarrow range
- status \rightarrow ignored(?)
- request \rightarrow returned from call

The arguments of various MPI calls get simplified:

- Communicator \rightarrow base object
- (pointer, size, type) \rightarrow range
- status \rightarrow ignored(?)
- request \rightarrow returned from call

Passing an owning range to non-blocking calls:

- Compiler error, or
- Payload gets bundled with the request

The arguments of various MPI calls get simplified:

- Communicator \rightarrow base object
- (pointer, size, type) \rightarrow range
- status \rightarrow ignored(?)
- request \rightarrow returned from call

Passing an owning range to non-blocking calls:

- Compiler error, or
- Payload gets bundled with the request

Error handling via exceptions or `std::expected`

Inputs:

- `std::vector` of in-neighbors (ranks)
- `std::flat_map` of out-neighbors (ranks) to corresponding message – vector of `doubles`

Inputs:

- `std::vector` of in-neighbors (ranks)
- `std::flat_map` of out-neighbors (ranks) to corresponding message – vector of `doubles`

Algorithm:

1. Post nonblocking sends
2. Probe for incoming message sizes
3. Gather incoming messages into single `std::vector`

```
1 std::vector<double> halo(MPI_Comm comm,
2                          const std::vector<int>& in_nbrs,
3                          const std::flat_map<int, std::vector<double>>& out) {
4     std::vector<MPI_Request> reqs;
5     send(comm, reqs, out);
6     const auto in_sizes = probe(comm, in_nbrs);
7     auto retval = receive(comm, reqs, in_nbrs, in_sizes);
8     MPI_Waitall(static_cast<int>(reqs.size()), reqs, MPI_STATUSES_IGNORE);
9     return retval;
10 }
```

```
1 std::vector<double> halo(MPI_Comm comm,
2                          const std::vector<int>& in_nbrs,
3                          const std::flat_map<int, std::vector<double>>& out) {
4     std::vector<MPI_Request> reqs;
5     send(comm, reqs, out);
6     const auto in_sizes = probe(comm, in_nbrs);
7     auto retval = receive(comm, reqs, in_nbrs, in_sizes);
8     MPI_Waitall(static_cast<int>(reqs.size()), reqs, MPI_STATUSES_IGNORE);
9     return retval;
10 }
```

```
1 std::vector<double> halo(const MPIComm& comm,
2                          const std::vector<int>& in_nbrs,
3                          const std::flat_map<int, std::vector<double>>& out) {
4     std::vector<MPIComm::Request> reqs;
5     send(comm, reqs, out);
6     const auto in_sizes = probe(comm, in_nbrs);
7     auto retval = receive(comm, reqs, in_nbrs, in_sizes);
8     // Waitall(reqs);
9     return retval;
10 }
```



```
1 void send(MPI_Comm comm,  
2           std::vector<MPI_Request>& reqs,  
3           const std::flat_map<int, std::vector<double>>& out) {  
4     for(const auto& [id, data] : out) {  
5       MPI_Request req;  
6       MPI_Isend(data.data(), data.size(), MPI_DOUBLE, id, 0, comm, &req);  
7       reqs.push_back(req);  
8     }  
9 }
```

```
1 void send(MPI_Comm comm,  
2           std::vector<MPI_Request>& reqs,  
3           const std::flat_map<int, std::vector<double>>& out) {  
4     for(const auto& [id, data] : out) {  
5       MPI_Request req;  
6       MPI_Isend(data.data(), data.size(), MPI_DOUBLE, id, 0, comm, &req);  
7       reqs.push_back(req);  
8     }  
9 }
```

```
1 void send(const MPIComm& comm,  
2           std::vector<MPIComm::Request>& reqs,  
3           const std::flat_map<int, std::vector<double>>& out) {  
4     for(const auto& [id, data] : out) {  
5       reqs.push_back(comm.Isend(data, id, 0));  
6     }  
7 }
```

```
1 void send(MPI_Comm comm,
2           std::vector<MPI_Request>& reqs,
3           const std::flat_map<int, std::vector<double>>& out) {
4     for(const auto& [id, data] : out) {
5         MPI_Request req;
6         MPI_Isend(data.data(), data.size(), MPI_DOUBLE, id, 0, comm, &req);
7         reqs.push_back(req);
8     }
9 }
```

```
1 void send(const MPIComm& comm,
2           std::vector<MPIComm::Request>& reqs,
3           const std::flat_map<int, std::vector<double>>& out) {
4     for(const auto& [id, data] : out) {
5         reqs.push_back(comm.Isend(data, id, 0));
6     }
7 }
```

```
1 void send(const MPIComm& comm,
2           std::vector<MPIComm::Request>& reqs,
3           const std::flat_map<int, std::vector<double>>& out) {
4     sr::transform(out, std::back_inserter(reqs), [&](const auto& pair){
5         return comm.Isend(pair.second, pair.first, 0);
6     });
7 }
```

```
1 std::vector<int> probe(MPI_Comm comm,
2                       const std::vector<int>& in_nbrs) {
3     std::vector<int> retval;
4     for(int id : in_nbrs) {
5         MPI_Status stat;
6         MPI_Probe(id, 0, comm, &stat);
7         MPI_Count sz;
8         MPI_Get_elements(&status, MPI_DOUBLE, &sz);
9         retval.push_back(sz);
10    }
11    return retval;
12 }
```

```
1 std::vector<int> probe(MPI_Comm      comm,
2                       const std::vector<int>& in_nbrs) {
3     std::vector<int> retval;
4     for(int id : in_nbrs) {
5         MPI_Status stat;
6         MPI_Probe(id, 0, comm, &stat);
7         MPI_Count sz;
8         MPI_Get_elements(&status, MPI_DOUBLE, &sz);
9         retval.push_back(sz);
10    }
11    return retval;
12 }
```

```
1 std::vector<int> probe(const MPIComm&      comm,
2                       const std::vector<int>& in_nbrs) {
3     return in_nbrs
4         | sv::transform([&](int id) { return comm.probeSize<double>(id, 0); })
5         | sr::to<std::vector>;
6 }
```

```
1 void receive(MPI_Comm      comm,
2              std::vector<MPI_Request>& reqs,
3              const std::vector<int>&   in_nbrs,
4              const std::vector<int>&   in_sizes) {
5     std::vector<double> retval(std::reduce(in_sizes.begin(), in_sizes.end()));
6     for(int i = 0, offset = 0; i != in_nbrs.size(); ++i) {
7         MPI_Request req;
8         MPI_Irecv(retval.data()+offset, in_sizes[i], MPI_DOUBLE, in_nbrs[i], 0, comm, &req);
9         reqs.push_back(req);
10        offset += sz;
11    }
12    return retval;
13 }
```

```
1 void receive(MPI_Comm          comm,
2              std::vector<MPI_Request>& reqs,
3              const std::vector<int>&   in_nbrs,
4              const std::vector<int>&   in_sizes) {
5     std::vector<double> retval(std::reduce(in_sizes.begin(), in_sizes.end()));
6     for(int i = 0, offset = 0; i != in_nbrs.size(); ++i) {
7         MPI_Request req;
8         MPI_Irecv(retval.data()+offset, in_sizes[i], MPI_DOUBLE, in_nbrs[i], 0, comm, &req);
9         reqs.push_back(req);
10        offset += sz;
11    }
12    return retval;
13 }
```

```
1 void receive(const MPIComm&      comm,
2              std::vector<MPIComm::Request>& reqs,
3              const std::vector<int>&   in_nbrs,
4              const std::vector<int>&   in_sizes) {
5     std::vector<double> retval(std::reduce(in_sizes.begin(), in_sizes.end()));
6     for(int offset = 0; const auto [id, sz] : sv::zip(in_nbrs, in_sizes)) {
7         reqs.push_back(comm.Irecv(retval | sv::drop(offset), id, 0));
8         offset += sz;
9     }
10    return retval;
11 }
```

Physics injection & Passing compile-time parameters

What we mean:

- Equations
- Boundary conditions
- Simulation structure, if impactful at compile-time
- Dimensions, if they meaningfully affect computation (e.g. stencils)

What we mean:

- Equations
- Boundary conditions
- Simulation structure, if impactful at compile-time
- Dimensions, if they meaningfully affect computation (e.g. stencils)

What we don't mean:

- Physical parameters, e.g., viscosity, heat capacity, ...
- Simulation structure, if no compile-time optimization occurs

2 distinct approaches:

2 distinct approaches:

- Code generation:
 - Flexibility
 - Requires auxiliary tooling
 - Hostile to IDEs
 - e.g.: TCLB

2 distinct approaches:

- Code generation:
 - Flexibility
 - Requires auxiliary tooling
 - Hostile to IDEs
 - e.g.: TCLB
- Meta-programming:
 - Self-contained
 - Transparent to the IDE
 - Requires some knowledge of C++ from the user
 - Templates
 - More templates!
 - Unless...?

We've always had aggregate initialization:

```
1 struct Params {  
2     int nx = 1, ny = 1, nz = 1;  
3     double nu = 1e-3;  
4     int n_rhs = 1, n_depends = 0;  
5 };  
6 Params p{2, 4, 1, 1e-3, 1, 2};
```

We've always had aggregate initialization:

```
1 struct Params {  
2     int nx = 1, ny = 1, nz = 1;  
3     double nu = 1e-3;  
4     int n_rhs = 1, n_depends = 0;  
5 };  
6 Params p{2, 4, 1, 1e-3, 1, 2};
```

But now we have *designated* initialization:

```
1 Params p{.nx = 2, .ny = 4, .n_depends = 2};
```

- Reader friendly
- No need to touch sensible defaults

We can now use any literal class type as a template parameter.
Combined with the following
`std::integral_constant`-esque utility

```
1 template<auto>  
2 struct ConstexprWrapper {};
```


We can now use any literal class type as a template parameter.
Combined with the following
`std::integral_constant`-esque utility

```
1 template<auto>  
2 struct ConstexprWrapper {};
```

it gives us a beautiful, clear, and powerful interface for passing
CT parameters:

```
1 constexpr auto params      = Params {2, 4, 1, 1e-3, 1, 2};  
2 /* 1 */   auto simulation = makeSim(ConstexprWrapper<params>{});  
3 /* 2 */   auto simulation = makeSim(WRAP_CT_PARAM(params));  
4 /* 3 */   auto simulation = makeSim<params>();
```

But we're not limited to simple structs:

```
1 // u, vort, p, grad T, T
2 constexpr auto sim_def = ProblemDef{
3     defineDomain< 11 >(IDfluid, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10),
4     defineDomain< 11 >(IDsolid, 7, 8, 9, 10)};
5 constexpr auto dbc_def = ProblemDef{
6     defineDomain< 11 >(wall, 7, 8, 9)};
7 auto sys = makeAlgebraicSystem<sim_def, dbc_def>(mesh /*,...*/);
```

example from L3STER

Since we know how to elegantly pass compile-time parameters, we know how to set the size of operators, the discretization order, etc.

```
1 constexpr KernelParams diff3_params{.dimension = 3,  
2                                     .n_equations = 7,  
3                                     .n_unknowns = 4};
```

example from L3STER

Since we know how to elegantly pass compile-time parameters, we know how to set the size of operators, the discretization order, etc.

```
1 constexpr KernelParams diff3_params{.dimension = 3,  
2                                     .n_equations = 7,  
3                                     .n_unknowns = 4};
```

example from L3STER

The question remains – how do we provide the bodies of these operators?

Since we know how to elegantly pass compile-time parameters, we know how to set the size of operators, the discretization order, etc.

```
1 constexpr KernelParams diff3_params{.dimension = 3,  
2                                     .n_equations = 7,  
3                                     .n_unknowns = 4};
```

example from L3STER

The question remains – how do we provide the bodies of these operators?

Idea: rely on generic lambdas and structured bindings to provide a type-oblivious interface.

```
1 const auto diffusion3d_kernel = wrapDomainEquationKernel<diff3_params>(
2     [](const auto& in, auto& out) {
3         auto& [operators, rhs] = out;
4         auto& [A0, Ax, Ay, Az] = operators;
5         constexpr double k = 1.; // diffusivity
6         constexpr double s = 1.; // source
7
8         Ax(0, 1) = -k;
9         Ay(0, 2) = -k;
10        Az(0, 3) = -k;
11        rhs[0] = s;
12        A0(1, 1) = -1.;
13        Ax(1, 0) = 1.;
14        A0(2, 2) = -1.;
15        Ay(2, 0) = 1.;
16        A0(3, 3) = -1.;
17        Az(3, 0) = 1.;
18        Ay(4, 3) = 1.;
19        Az(4, 2) = -1.;
20        Ax(5, 3) = -1.;
21        Az(5, 1) = 1.;
22        Ax(6, 2) = 1.;
23        Ay(6, 1) = -1.;
24    });
25 sys->assembleProblem(diffusion3d_kernel, std::views::single(solid_id));
```

example from L3STER

- Today's C++ is not the same language you held your nose at 20 years ago

- Today's C++ is not the same language you held your nose at 20 years ago
- It's (relatively) easy to pass compile-time information to an API and to the compiler

- Today's C++ is not the same language you held your nose at 20 years ago
- It's (relatively) easy to pass compile-time information to an API and to the compiler
- Most of the complexity can be hidden from the user

- Today's C++ is not the same language you held your nose at 20 years ago
- It's (relatively) easy to pass compile-time information to an API and to the compiler
- Most of the complexity can be hidden from the user

Discussion

`jakub.galecki.dokt@pw.edu.pl`